# **SAMPLE OUTPUT**

# **HARP**

# 12 Threads, 370K Records

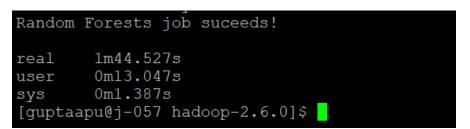
[guptaapu@j-057 hadoop-2.6.0]\$ time hadoop jar contrib-0.1.0.jar edu.iu.rf.RFMapCollective 32 2 12 /data/train\_test\_files/train\_airline\_harp\_5.csv /data/train\_test\_files/merged\_test\_airline\_harp\_5.csv /out\_32\_2\_12\_harp

Get map collective runner



## Collecting Memory Consumption data

Dec 02 2018 16:39:36 The Intel VTune Amplifier 2019 is collecting data. Closing this window will cancel the analysis and terminate the profiled application in case the Launch Application target type is selected.





#### ▼ Top Memory-Consuming Functions

Total Thread Count: 63
Paused Time <sup>3</sup>: 0s

This section lists the most memory-consuming functions in your application.

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
os::malloc	39 MB	8 MB	2,847	libjvm.so
init_tpssmrte_jstring	19 MB	896 KB	10,086	libamplxe_tpssmrte_javaprof_1.0.so
os::Linux::commit_memory_impl	1 MB	1 MB	31	libjvm.so
Interpreter	1 MB	193 KB	704	[Dynamic code]
java::lang::ClassLoader::defineClass1	867 KB	557 KB	2,523	[Compiled Java code]
[Others]	1 MB	152 KB	32,256	

\*N/A is applied to non-summable metrics.

### ○ Collection and Platform Info

This section provides information about this collection, including result set size and collection platform data.

Application Command Line:

Operating System:

3.10.0-862.9.1.eI7.x86\_64 NAME="Red Hat Enterprise Linux Server" VERSION="7.5 (Maipo)" ID="rhel" ID\_LIKE="fedora" VARIANT="Server" VARIANT\_ID="server" VERSION\_ID="7.5" PRETTY\_NAME="Red Hat Enterprise Linux Server 7.5 (Maipo) ANSI\_COLOR="0;31" CPE\_NAME="cpe:/o:redhat:enterprise\_linux:7.5:GA:server" HOME\_URL="https://www.redhat.com/" BUG\_REPORT\_URL="https://bugzilla.redhat.com/" REDHAT\_BUGZILLA\_PRODUCT="Red Hat Enterprise Linux 7" REDHAT\_BUGZILLA\_PRODUCT\_VERSION=7.5 REDHAT\_SUPPORT\_PRODUCT="Red Hat Enterprise Linux"

REDHAT\_SUPPORT\_PRODUCT\_VERSION="7.5"

Computer Name: j-057 Result Size: 65 MB

Collection start time: 21:39:36 02/12/2018 UTC Collection stop time: 21:41:04 02/12/2018 UTC Collector Type: User-mode sampling and tracing

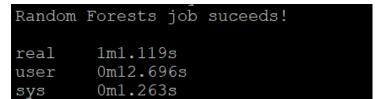
Finalization mode: Fast. If the number of collected samples exceeds the threshold, this mode limits the number of processed samples to speed up post-processing.

Intel(R) Xeon(R) E5/E7 v3 Processor code named Haswell

Frequency: 2.3 GHz Logical CPU Count: 48

## 24 Threads, 370K Records

guptaapu@j-057 hadoop-2.6.0]\$ time hadoop jar contrib-0.1.0.jar edu.iu.rf.RFMapCollective 32 2 24 /data/train\_test\_files/train\_airline\_harp\_5.csv /data/trai \_test\_files/merged\_test\_airline\_harp\_5.csv /out\_32\_2\_24\_harp





# 

Allocation Size: 40 MB Deallocation Size: 32 MB Allocations: 23,609 Total Thread Count: 62 Paused Time 2:

## ▼ Top Memory-Consuming Functions

This section lists the most memory-consuming functions in your application.

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
os::malloc	26 MB	6 MB	1,533	libjvm.so
init_tpssmrte_jstring	9 MB	356 KB	4,766	libamplxe_tpssmrte_javaprof_1.0.so
Interpreter	1 MB	310 KB	554	[Dynamic code]
java::lang::ClassLoader::defineClass1	835 KB	525 KB	2,522	[Compiled Java code]
os::Linux::commit_memory_impl	544 KB	520 KB	12	libjvm.so
[Others]	1002 KB	105 KB	14,222	

<sup>\*</sup>N/A is applied to non-summable metrics.

### ○ Collection and Platform Info

This section provides information about this collection, including result set size and collection platform data.

Application Command Line:

3.10.0-862.9.1.eI7.x86 64 NAME="Red Hat Enterprise Linux Server" VERSION="7.5 (Maipo)" ID="rhel" ID LIKE="fedora" Operating System:

VARIANT="Server" VARIANT\_ID="server" VERSION\_ID="7.5" PRETTY\_NAME="Red Hat Enterprise Linux Server 7.5 (Maipo)" ANSI\_COLOR="0;31" CPE\_NAME="cpe:/o:redhat:enterprise\_linux:7.5:GA:server" HOME\_URL="https://www.redhat.com/" BUG\_REPORT\_URL="https://bugzilla.redhat.com/" REDHAT\_BUGZILLA\_PRODUCT="Red Hat Enterprise Linux 7" REDHAT\_BUGZILLA\_PRODUCT\_VERSION=7.5 REDHAT\_SUPPORT\_PRODUCT="Red Hat Enterprise Linux" REDHAT\_SUPPORT\_PRODUCT\_VERSION="7.5"

Computer Name: j-057 56 MB Result Size:

Collection start time: 21:43:51 02/12/2018 UTC Collection stop time: 21:44:30 02/12/2018 UTC Collector Type: User-mode sampling and tracing

Finalization mode: Fast. If the number of collected samples exceeds the threshold, this mode limits the number of processed samples to speed up post-processing.

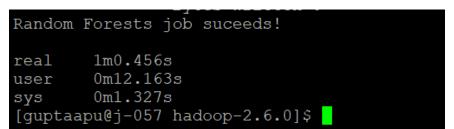
Name:

Intel(R) Xeon(R) E5/E7 v3 Processor code named Haswell

Frequency: Logical CPU Count: 48

## 48 Threads, 370 K Records

\_ guptaapu@j-057 hadoop-2.6.0]\$ time hadoop jar contrib-0.1.0.jar edu.iu.rf.RFMapCollective 32 2 48 /data/train\_test\_files/train\_airline\_harp\_5.csv /data/trai \_test\_files/merged\_test\_airline\_harp\_5.csv /out\_32\_2\_48\_harp t map collective runner





## ▼ Top Memory-Consuming Functions

This section lists the most memory-consuming functions in your application.

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
os::malloc	28 MB	6 MB	1,617	libjvm.so
init_tpssmrte_jstring	9 MB	384 KB	5,012	libamplxe_tpssmrte_javaprof_1.0.so
Interpreter	1 MB	343 KB	571	[Dynamic code]
os::Linux::commit_memory_impl	864 KB	840 KB	14	libjvm.so
java::lang::ClassLoader::defineClass1	803 KB	493 KB	2,521	[Compiled Java code]
[Others]	1007 KB	169 KB	15,052	

\*N/A is applied to non-summable metrics

#### Collection and Platform Info

This section provides information about this collection, including result set size and collection platform data.

3.10.0-862.9.1.el7.x86\_64 NAME="Red Hat Enterprise Linux Server" VERSION="7.5 (Maipo)" ID="rhel" ID\_LIKE="fedora" Operating System:

VARIANT="Server" VARIANT\_ID="server" VERSION\_ID="7.5" PRETTY\_NAME="Red Hat Enterprise Linux Server 7.5 (Maipo)" VARIANT = Server VARIANT \_ID = Server VERSION\_ID = 7.5 PRETIT\_NAME = Red Rat Enterprise Linux Server 7.5 (Ma ANSI\_COLOR="0;31" CPE\_NAME="cpe\_/o:redhat.enterprise\_linux:7.5:GA:server" HOME\_URL="https://www.redhat.com/" BUG\_REPORT\_URL="https://www.redhat.com/" REDHAT\_BUGZILLA\_PRODUCT="Red Hat Enterprise Linux 7" REDHAT\_BUGZILLA\_PRODUCT\_VERSION=7.5 REDHAT\_SUPPORT\_PRODUCT="Red Hat Enterprise Linux" REDHAT\_SUPPORT\_PRODUCT\_VERSION="7.5"

Computer Name: j-057 Result Size: 56 MB

Collection start time: 21:46:30 02/12/2018 UTC Collection stop time: 21:47:15 02/12/2018 UTC Collector Type: User-mode sampling and tracing

Finalization mode: Fast. If the number of collected samples exceeds the threshold, this mode limits the number of processed samples to speed up post-processing.

Name: Intel(R) Xeon(R) E5/E7 v3 Processor code named Haswell

Frequency: Logical CPU Count: 48

# **SPARK**

## 12 Threads, 370 K Records

hadoop2.7]\$ time bin/spark-submit rf\_spark.py /data/train\_test\_files/train\_airline\_spark\_5.csv /data/train\_test\_files/ NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable SparkContext:54 - Running Spark version 2.3.2 SparkContext:54 - Submitted application: AirlineApp

```
Building Random Forest Model..
Measuring test results:
Test Error = 0.0699952
f1 = 0.927965
        0m25.646s
real
        2m42.103s
user
        0m14.374s
```

**INTEL VTUNE AMPLIFIER 2019** Memory Consumption Memory Consumption 🔻 🕐 Analysis Configuration Collection Log Summary Bottom-up

## 

Allocation Size: 2 MB Deallocation Size: 2 MB 2,864 Total Thread Count: 2 Paused Time 2:

## ▼ Top Memory-Consuming Functions

This section lists the most memory-consuming functions in your application.

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
sock_recv	2 MB	0 B	645	_socket.so
сору	96 KB	0 B	24	copy.py
init	71 KB	9 KB	95	java_gateway.py
init	69 KB	3 KB	92	java_gateway.py
IO_StringIO	40 KB	128 B	322	cStringIO.so
[Others]	178 KB	288 B	1,686	

\*N/A is applied to non-summable metrics

### 48 Threads, 370 K Records

```
rk-2.3.2-bin-hadoop2.7]$ time bin/spark-submit rf_spark.py /data/train_test_files/train_airline_spark_5.csv /data/train_test_files/test
line spark 5.csv 32
                           NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
   -12-02 16:51:02 WARN
                           SparkContext:54 - Running Spark version 2.3.2
SparkContext:54 - Submitted application: AirlineApp
```

INTEL VTUNE AMPLIFIER 2019

```
Building Random Forest Model..
Measuring test results:
Test Error = 0.0699952
f1 = 0.927965
        0m27.808s
real
        3m0.972s
user
        0m14.865s
sys
```

# Analysis Configuration Collection Log Summary Bottom-up ✓ Elapsed Time <sup>②</sup>: 10.645s

Allocation Size: 2 MB Deallocation Size: 2 MB Allocations: 2,869

Total Thread Count: 2 Paused Time <sup>②</sup>: 0s

## ▼ Top Memory-Consuming Functions

Memory Consumption Memory Consumption 🔻 🕜

This section lists the most memory-consuming functions in your application.

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
sock_recv	2 MB	0 B	645	_socket.so
сору	96 KB	0 B	24	copy.py
init	71 KB	9 KB	95	java_gateway.py
init	69 KB	3 KB	92	java_gateway.py
IO_StringIO	40 KB	128 B	322	cStringIO.so
[Others]	190 KB	12 KB	1,691	

\*N/A is applied to non-summable metrics.

#### Collection and Platform Info

This section provides information about this collection, including result set size and collection platform data.

Application Command Line: Operating System:

3.10.0-862.9.1.el7.x86\_64 NAME="Red Hat Enterprise Linux Server" VERSION="7.5 (Maipo)" ID="rhel" ID\_LIKE="fedora" VARIANT="Server" VARIANT\_ID="server" VERSION\_ID="7.5" PRETTY\_NAME="Red Hat Enterprise Linux Server 7.5 (Maipo)"

ANSI\_COLOR="0;31" CPE\_NAME="ope/o:redhat.enterprise\_linux.75:PHELITY\_NAME="Hed Hat Enterprise Linux Server 7.5 (Ma ANSI\_COLOR="0;31" CPE\_NAME="ope/o:redhat.com/" BLG\_REPORT\_URL="https://www.redhat.com/" REDHAT\_BUGZILLA\_PRODUCT="Red Hat Enterprise Linux 7" REDHAT\_BUGZILLA\_PRODUCT\_VERSION=7.5 REDHAT\_SUPPORT\_PRODUCT="Red Hat Enterprise Linux" REDHAT\_SUPPORT\_PRODUCT\_VERSION="7.5"

Computer Name: j-057 Result Size: 4 MB

Collection start time: 21:51:17 02/12/2018 UTC 21:51:28 02/12/2018 UTC Collection stop time: Collector Type: User-mode sampling and tracing

Finalization mode: Fast. If the number of collected samples exceeds the threshold, this mode limits the number of processed samples to speed up post-processing.

Name: Intel(R) Xeon(R) E5/E7 v3 Processor code named Haswell

Frequency: Logical CPU Count: 48